

# Workout Classification using LeNet and AlexNet

Alessia Marzotti  
Anuja Saira Abraham



# Our original dataset

Our dataset contains images from 22 different classes of workout exercises:

Examples: bench press, pull up, push up, leg raises, hip thrust,...

**Project Goal:** build a model which is able to predict the workout exercise class given an image

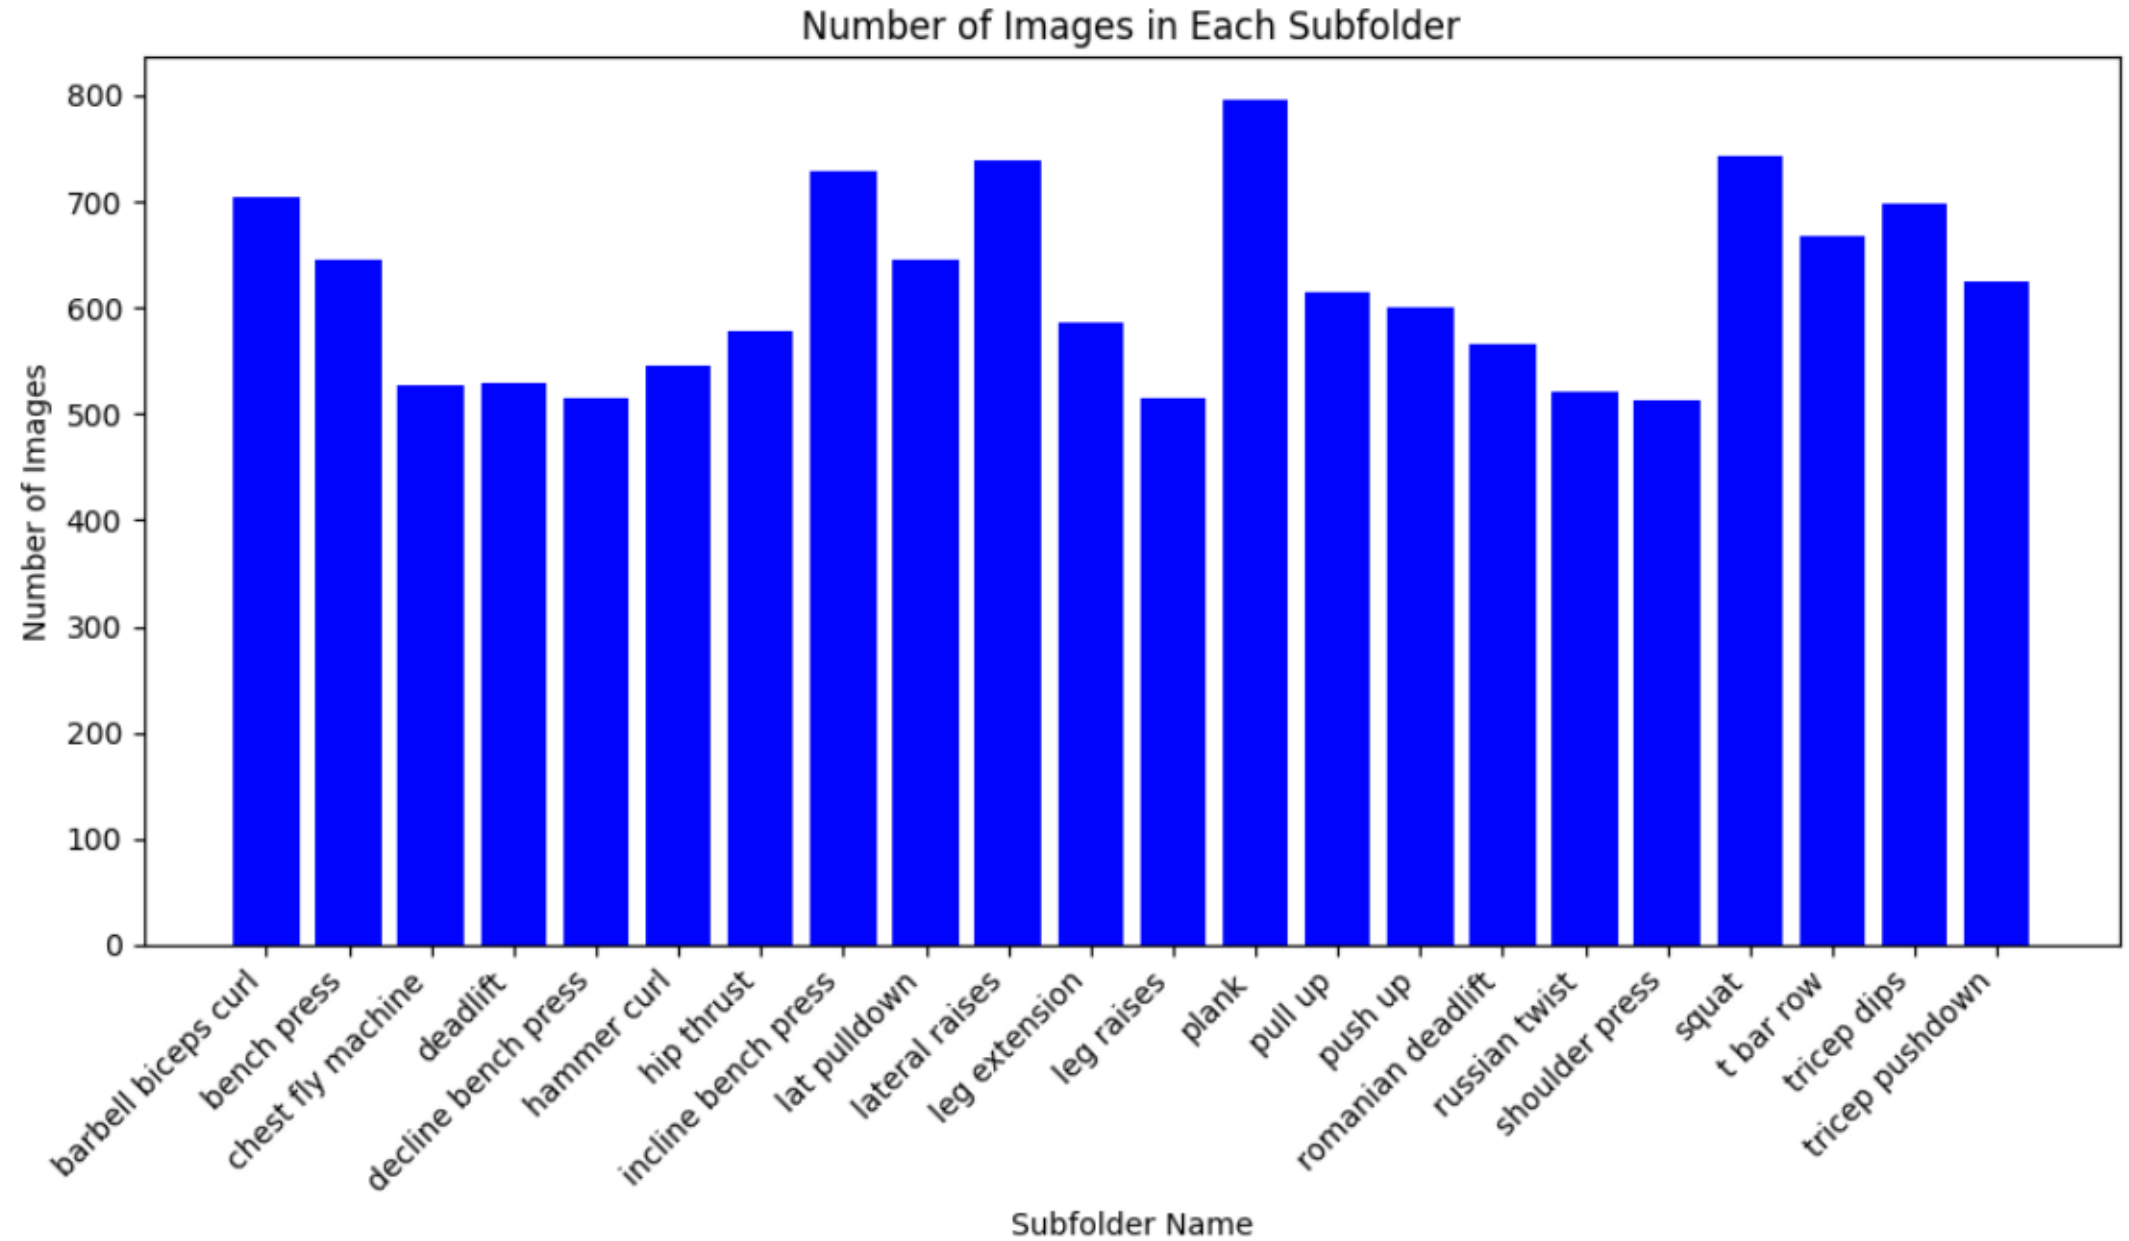
*Bench press*



*Pull up*



Our final dataset:



# Data Preparation

1

## **Data balancing:**

We have rebalanced the dataset because some classes were having too many images compared to other classes.

2

## **Data splitting:**

We have randomly split the data into train, validation and test (80%-10%-10%).

3

## **Data augmentation:**

Methods implemented:

- Resize
- RandomHorizontalFlip
- Grayscale

# First Architecture


## LeNet

Our **implementation** is composed by:

- A 11x11 convolutional layer, followed by an average pooling
- A 5X5 convolutional layer, followed by an average pooling
- A flatten layer
- A fully connected layer 1
- A fully connected layer 2
- A dropout layer
- A fully connected layer 3

**Model result:**

loss 0.642, train acc 0.823, val acc 0.851  
1071.7 examples/sec on cuda:0



# Implementation:

```
net = torch.nn.Sequential(  
    nn.Conv2d(1, 6, kernel_size=11, padding=2), nn.ReLU(),  
    nn.AvgPool2d(kernel_size=2, stride=2),  
    nn.Conv2d(6, 16, kernel_size=5), nn.ReLU(),  
    nn.AvgPool2d(kernel_size=2, stride=2),  
    nn.Flatten(),  
    nn.Linear(33856, 120), nn.ReLU(),  
    nn.Linear(120, 84), nn.ReLU(),  
    nn.Dropout(p=0.6),  
    nn.Linear(84, 22)  
)
```

## *Notes:*

- We applied a bigger kernel size
- We used a ReLU activation function
- We inserted a Dropout layer

# LeNet Predictions

Some examples:

True: tricep dips  
Predicted: tricep dips



True: plank  
Predicted: plank



True: squat  
Predicted: squat



True: tricep dips  
Predicted: squat




True: tricep dips  
Predicted: lat pulldown



We also created a function to *display misclassified* images:

The overall test accuracy is 83.8%





# Second architecture: AlexNet


The architecture is composed by:

- A 11x11 convolution layer 1, followed by a max pooling
- A 5x5 convolutional layer 2, followed by a max pooling
- A 3x3 convolutional layer 3
- A 3x3 convolutional layer 4
- A 3x3 convolutional layer 5, followed by a max pooling
- A flatten layer
- A fully connected layer 1
- A Dropout layer 1
- A fully connected layer 2
- A dropout layer 2
- A fully connected layer 3

```
net = nn.Sequential(  
    nn.Conv2d(1, 96, kernel_size=11, stride=4, padding=1), nn.ReLU(),  
    nn.MaxPool2d(kernel_size=3, stride=2),  
    nn.Conv2d(96, 256, kernel_size=5, padding=2), nn.ReLU(),  
    nn.MaxPool2d(kernel_size=3, stride=2),  
    nn.Conv2d(256, 384, kernel_size=3, padding=1), nn.ReLU(),  
    nn.Conv2d(384, 384, kernel_size=3, padding=1), nn.ReLU(),  
    nn.Conv2d(384, 256, kernel_size=3, padding=1), nn.ReLU(),  
    nn.MaxPool2d(kernel_size=3, stride=2),  
    nn.Flatten(),  
    nn.Linear(6400, 4096), nn.ReLU(),  
    nn.Dropout(p=0.5),  
    nn.Linear(4096, 4096), nn.ReLU(),  
    nn.Dropout(p=0.5),  
    nn.Linear(4096, 22))
```

loss 0.196, train acc 0.944, val acc 0.951  
721.9 examples/sec on cuda:0

**Model result:**



# AlexNet Predictions

Some examples:

True: pull up  
Predicted: pull up



True: romanian deadlift  
Predicted: romanian deadlift




True: bench press  
Predicted: bench press



True: barbell biceps curl  
Predicted: lateral raises




We created a function to display  
*missclassified images*:



The overall  
Test Accuracy  
Is 95.39%

# Conclusions

	LeNet	AlexNet
Epochs	9	9
Training loss	0.642	0.196
Training accuracy	82.3%	94.4%
Test accuracy	83.8%	95.39%
Examples per seconds	1071.7	721.9



Thank you for your  
attention!